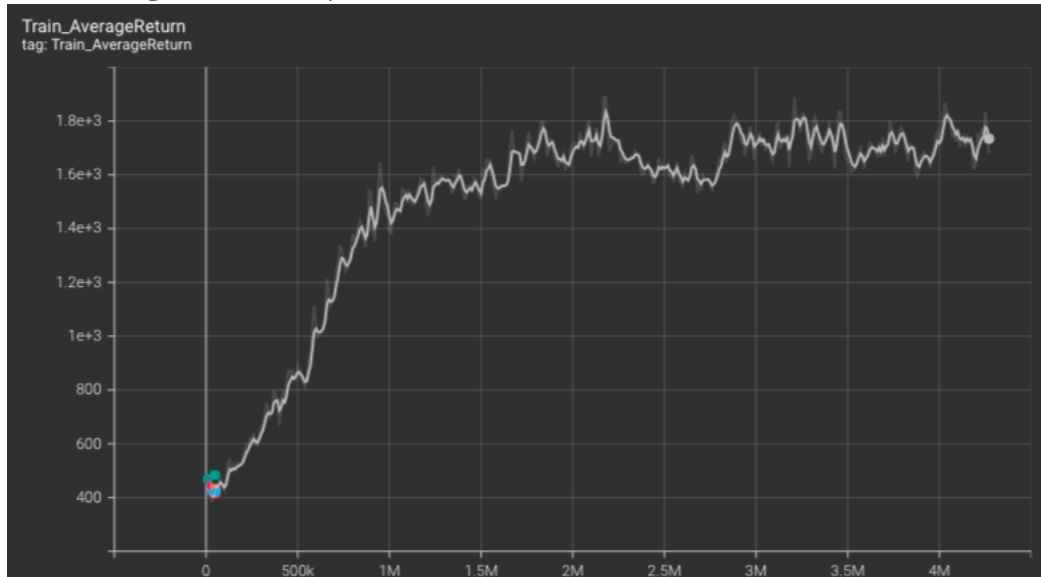



# Homework #3 Report

Alejandro Escontrela

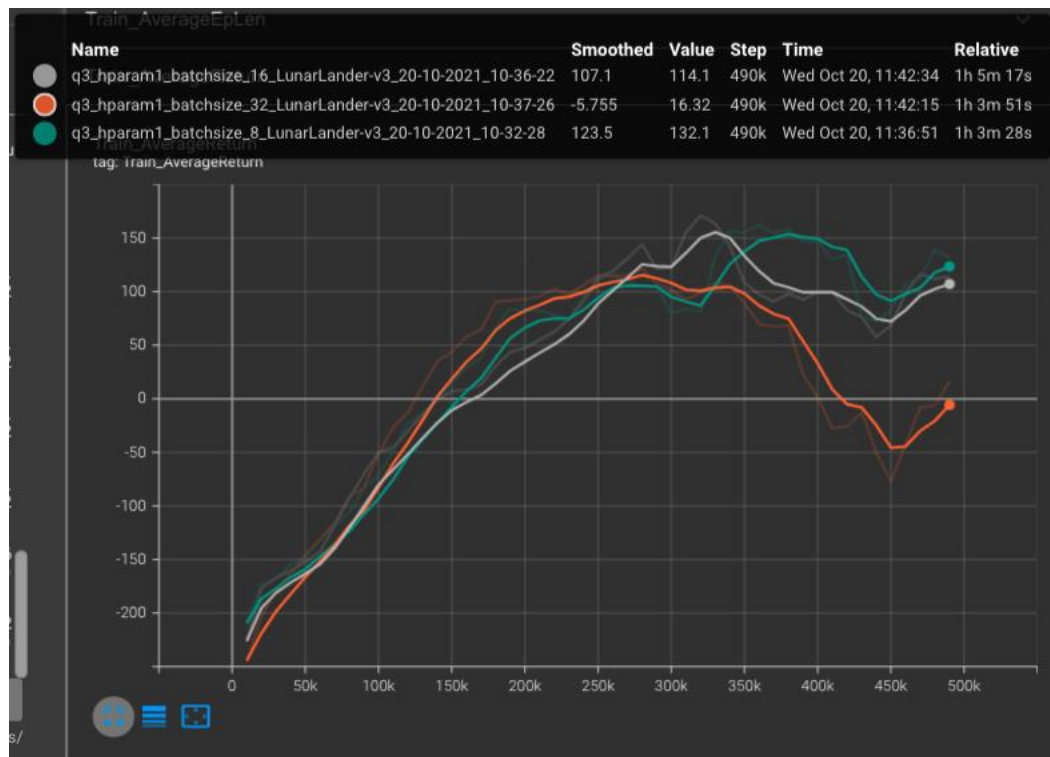
## Q1: Average return on pacman



## Q2: Double Q-Learning (DDQN)

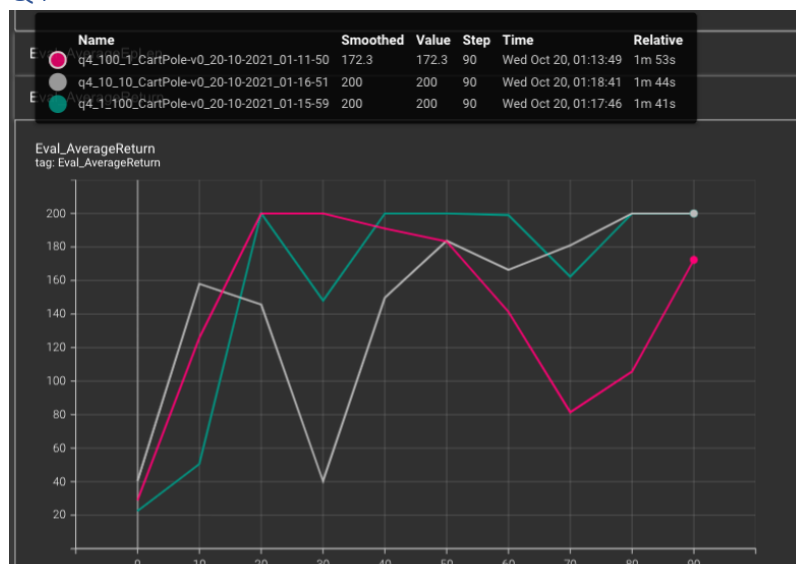


## Q3: DQN Hyperparameter Sweep




For this part of the assignment, I performed a sweep over the batchsize. Using a smaller batchsize leads to more stable learning.

#### Q4



While all three configurations achieve the highest score, it is evident that using 100 gradient steps and one critic update per update leads to more stable performance.

## Q5 – Inverted Pendulum



## Q5 – HalfCheetah

